
Zero-knowledge proofs security, in practice



JP Aumasson
@veorq

CSO @ taurushq.com

NIZK arguments security, in practice



JP Aumasson
@veorq

CSO @ taurushq.com

This talk

Focus on **zkSNARKs**, a class of NIZK arguments

- *Fully succinct* = $O(1)$ proof size and $O(\text{circuit size})$ verification time

Based on my experience looking for bugs in systems using

- **Groth16**, used in Zcash, Filecoin, and many others
- **Marlin**, a universal zkSNARK, used in Aleo

Most of the content applies to other systems (Plonk, SONIC, etc.) and STARKs

Why study zkSNARKs security?

For blockchain projects: **A major risk**:

- Complexity + Novelty => Non-trivial **bugs**
- A lot **at stake** (\$\$\$ and user data/privacy)

Why study zkSNARKs security?

For blockchain projects: **A major risk**:

- Complexity + Novelty => Non-trivial bugs
- A lot **at stake** (\$\$\$ and user data/privacy)

As a cryptographer: **The most interesting** crypto today:

- Solving real-world problems and deployed at scale (good papers + good code!)
- Intricate constructions with non-trivial components
- “Simple but complex” (non-interactive, many moving parts)
- “Multidimensional” way to reason about security

What is zkSNARKs security?

Soundness, often the *highest risk* in practice:

- Invalid proofs should always be rejected
- Forging, altering, replaying valid proofs should be impossible

What is zkSNARKs security?

Soundness, often the *highest risk* in practice:

- Invalid proofs should always be rejected
- Forging, altering, replaying valid proofs should be impossible

Zero-knowledge: Proofs should not leak witness information (private variables)

- In practice succinct proofs of large programs can leak only little data

What is zkSNARKs security?

Soundness, often the *highest risk* in practice:

- Invalid proofs should always be rejected
- Forging, altering, replaying valid proofs should be impossible

Zero-knowledge: Proofs should not leak witness information (private variables)

- In practice succinct proofs of large programs can leak only little data

Completeness, often a DoS/usability risk that may be further exploited:

- Valid proofs should always be accepted
- All programs/circuits supported should be correctly processed

Bug hunting challenges

Practical zkSNARKs are recent, thus auditors often have

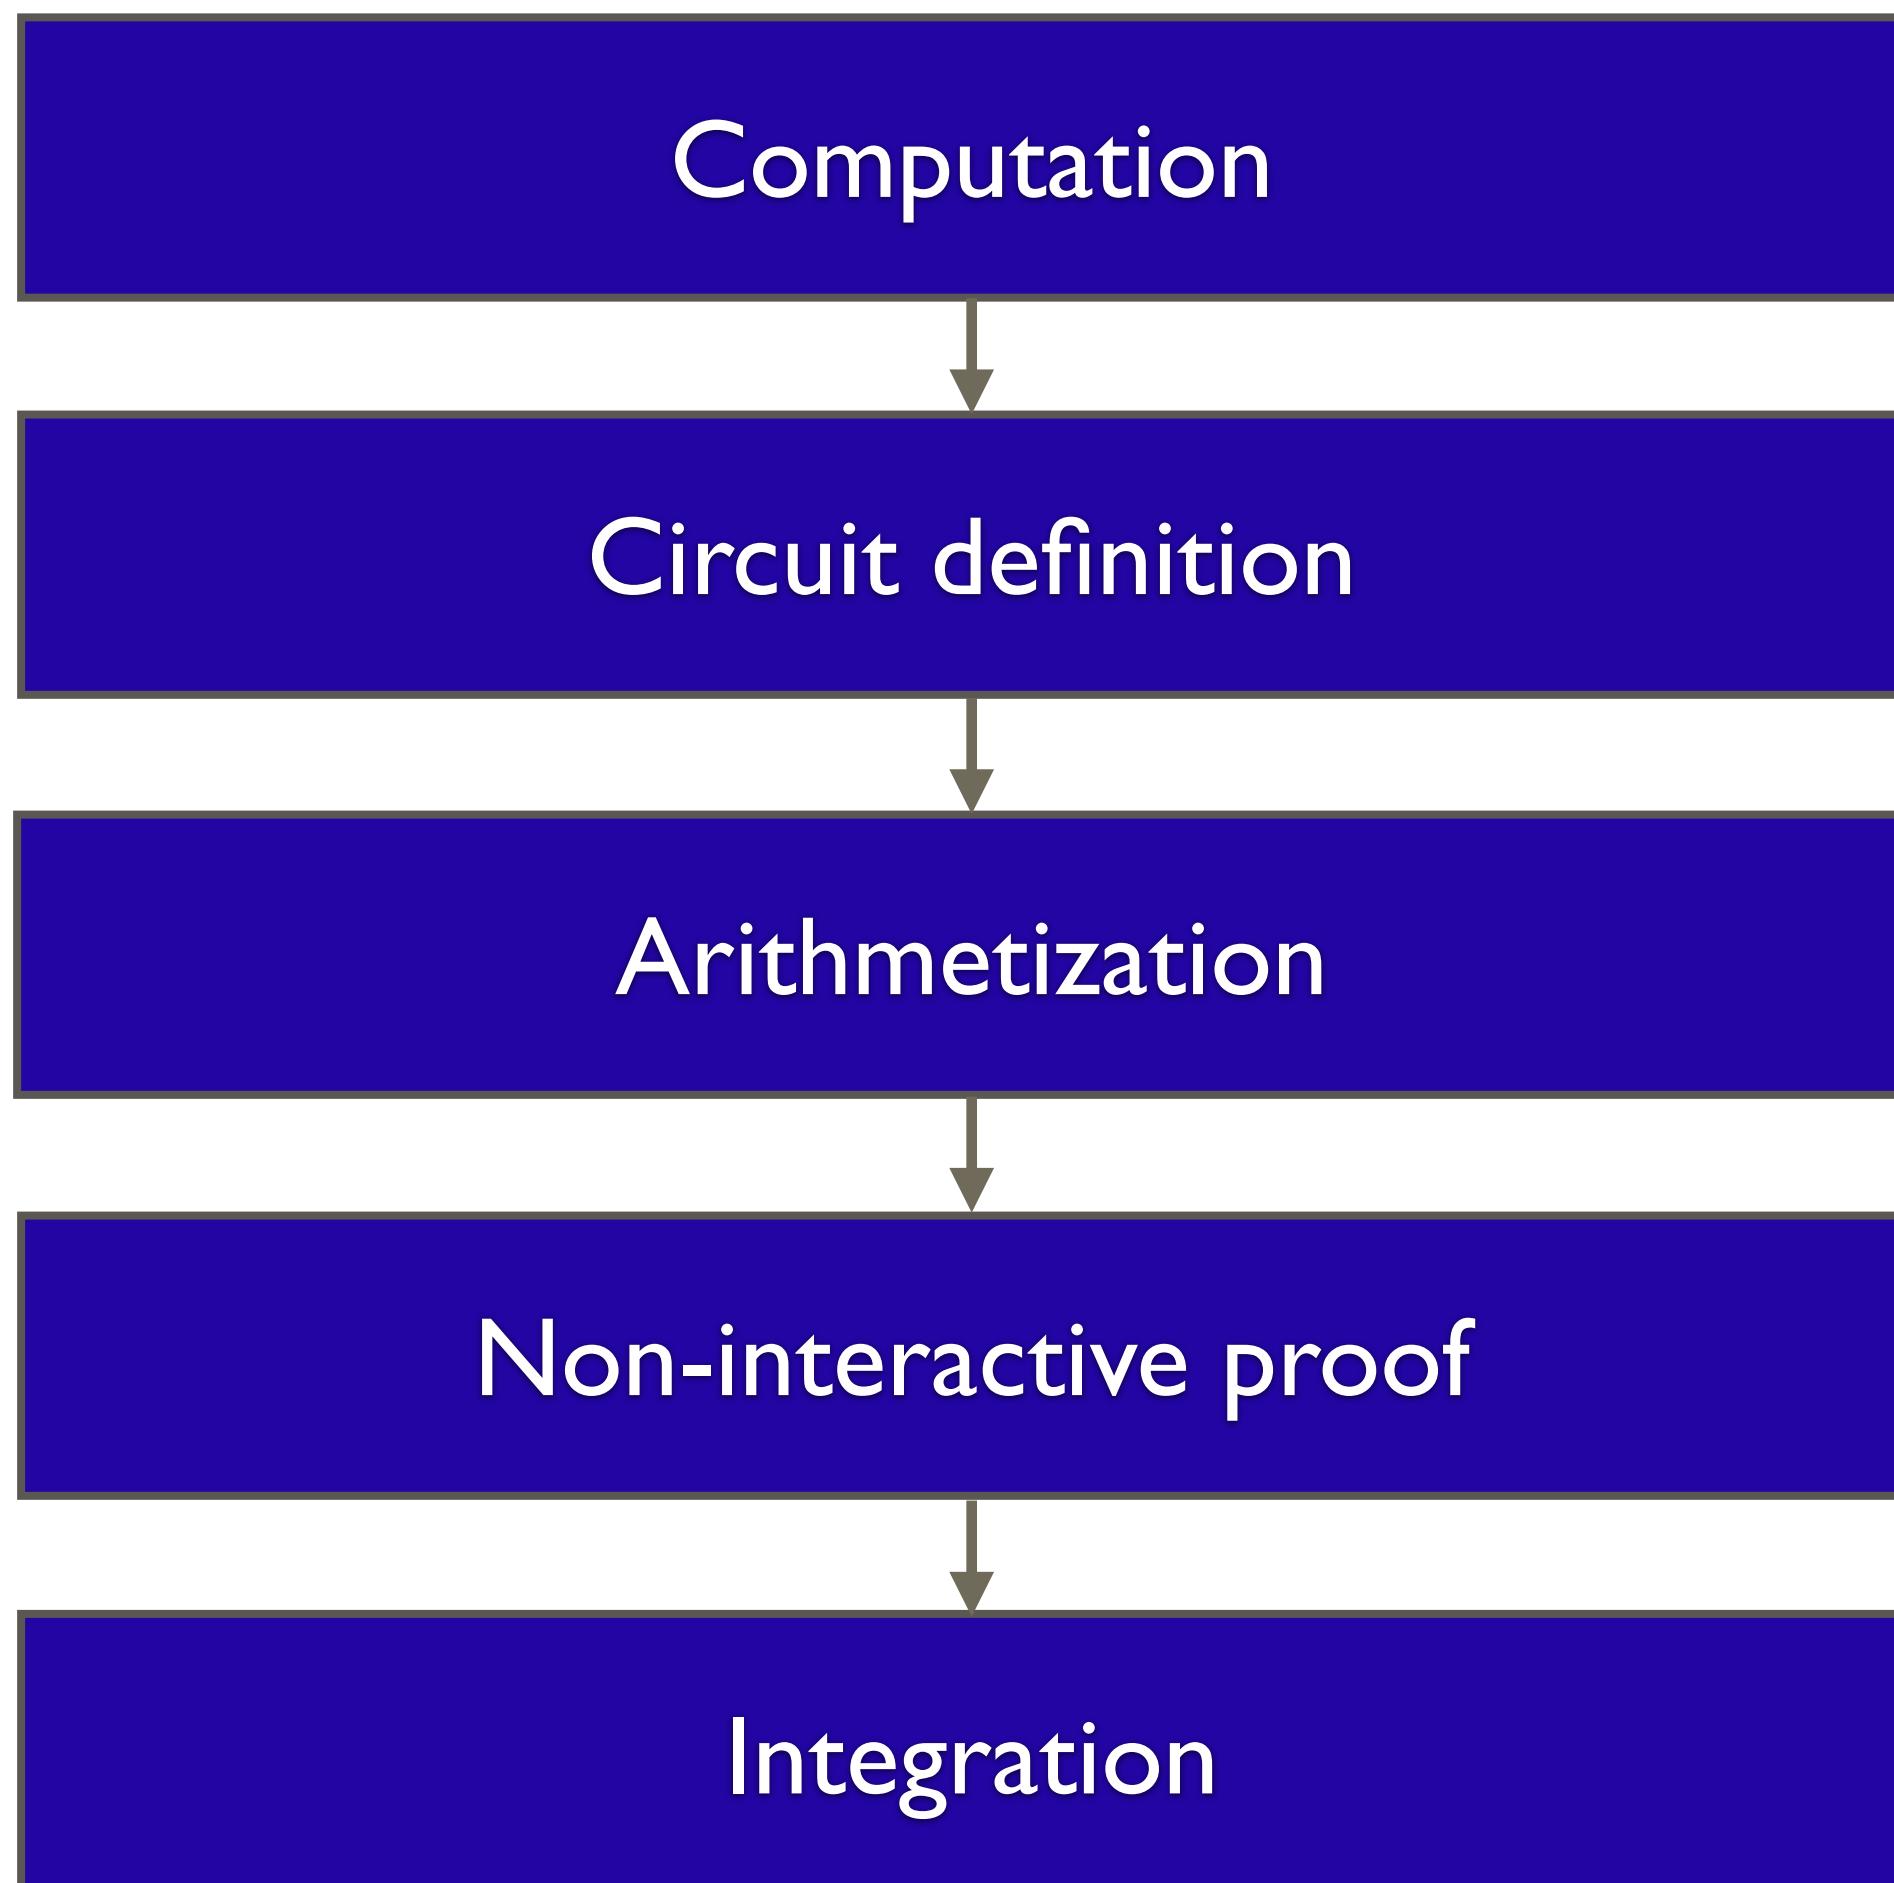
- Limited **experience** auditing ZKPs
- Limited **knowledge** of the theory, of implementations' tricks
- Limited "**checklist**" of bugs and bug classes
- Limited **tooling** and methodology

Most bugs found internally or by teams of similar projects

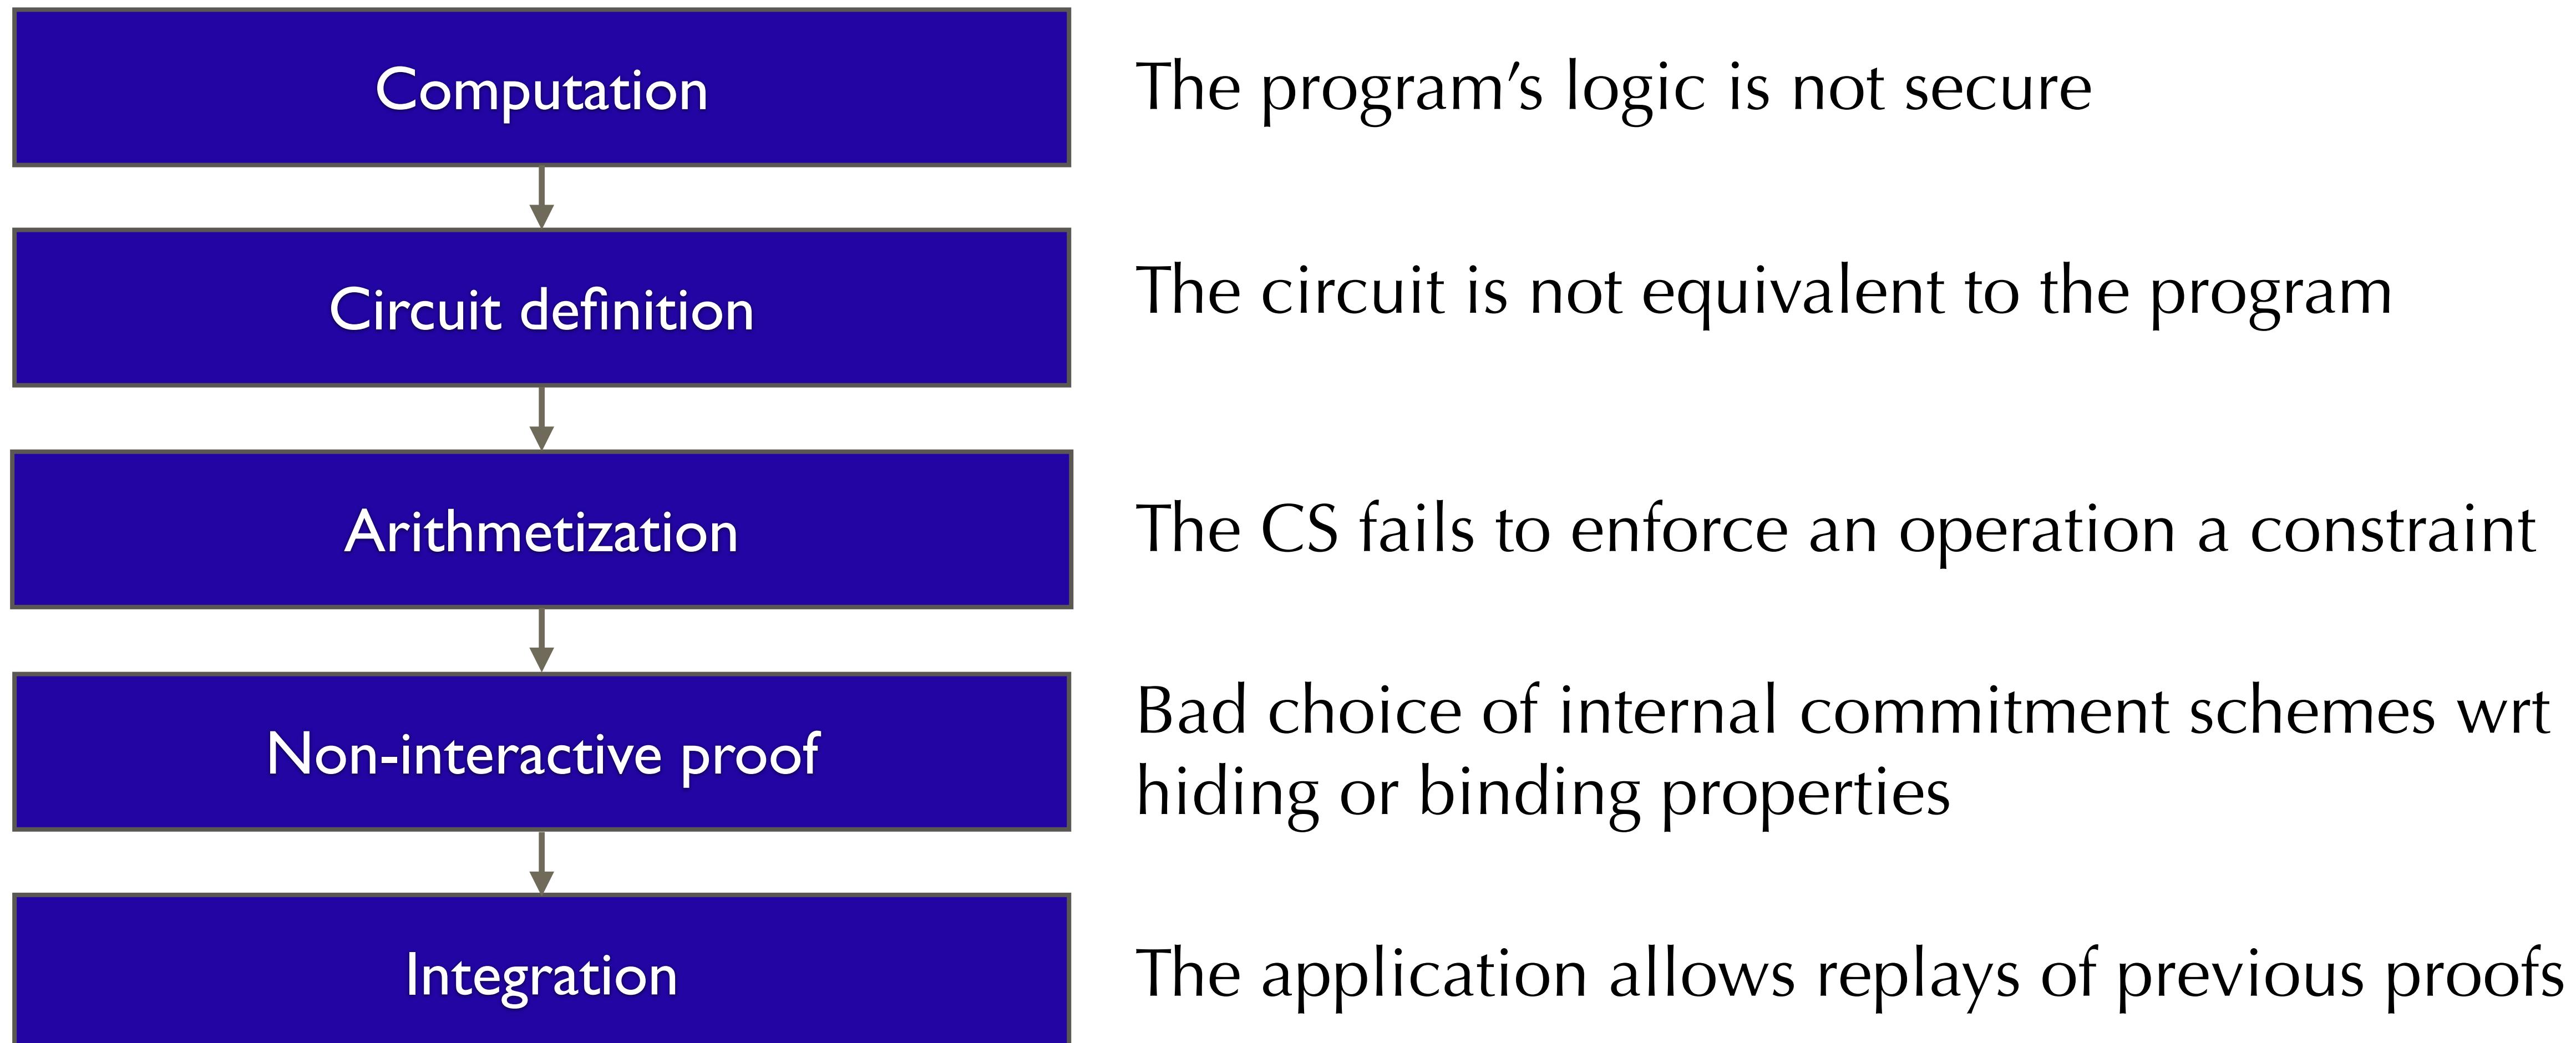
New crypto, new approach

- More **collaboration** with the devs/designers (joint review sessions, Q&As, etc.)
- More **threat analysis**, to understand the application's unique/novel risks
- Practical **experience**: writing PoCs, circuits, proof systems, etc.
- Learn **previous failures**, for example from...
 - Public disclosures and exploits
 - Other audit reports
 - Issue trackers / PRs
 - Community

General workflow, and failure examples



General workflow, and failure examples



How to break zkSNARKs security? (1/2)

Break soundness, for example by exploiting

- Constraint system not effectively enforcing certain constraints
- Insecure generation or protection of proving keys

How to break zkSNARKs security? (1/2)

Break soundness, for example by exploiting

- Constraint system not effectively enforcing certain constraints
- Insecure generation or protection of proving keys

Break zero-knowledge, for example by exploiting

- Private data treated as public variables
- Protocol-level “metadata attacks”

How to break zkSNARKs security? (1/2)

Break soundness, for example by exploiting

- Constraint system not effectively enforcing certain constraints
- Insecure generation or protection of proving keys

Break zero-knowledge, for example by exploiting

- Private data treated as public variables
- Protocol-level “metadata attacks”

Break completeness, for example by exploiting

- Incorrect R1CS synthesis behaviour on edge cases (e.g. wrt number of private vars)
- Gadget composition failure caused by type mismatch between gadget i/o values

How to break zkSNARKs security? (2/2)

Break (off-chain) software, via any bug leading to

- Leakage of data, including via side channels, encodings (“ZK execution”)
- Any form in insecure state (code execution, DoS)

Compromise the supply-chain, via

- Trusted setup's code and execution
- Build and release process integrity
- Software dependencies

How to break zkSNARKs security? (2/2)

Break (off-chain) software, via any bug leading to

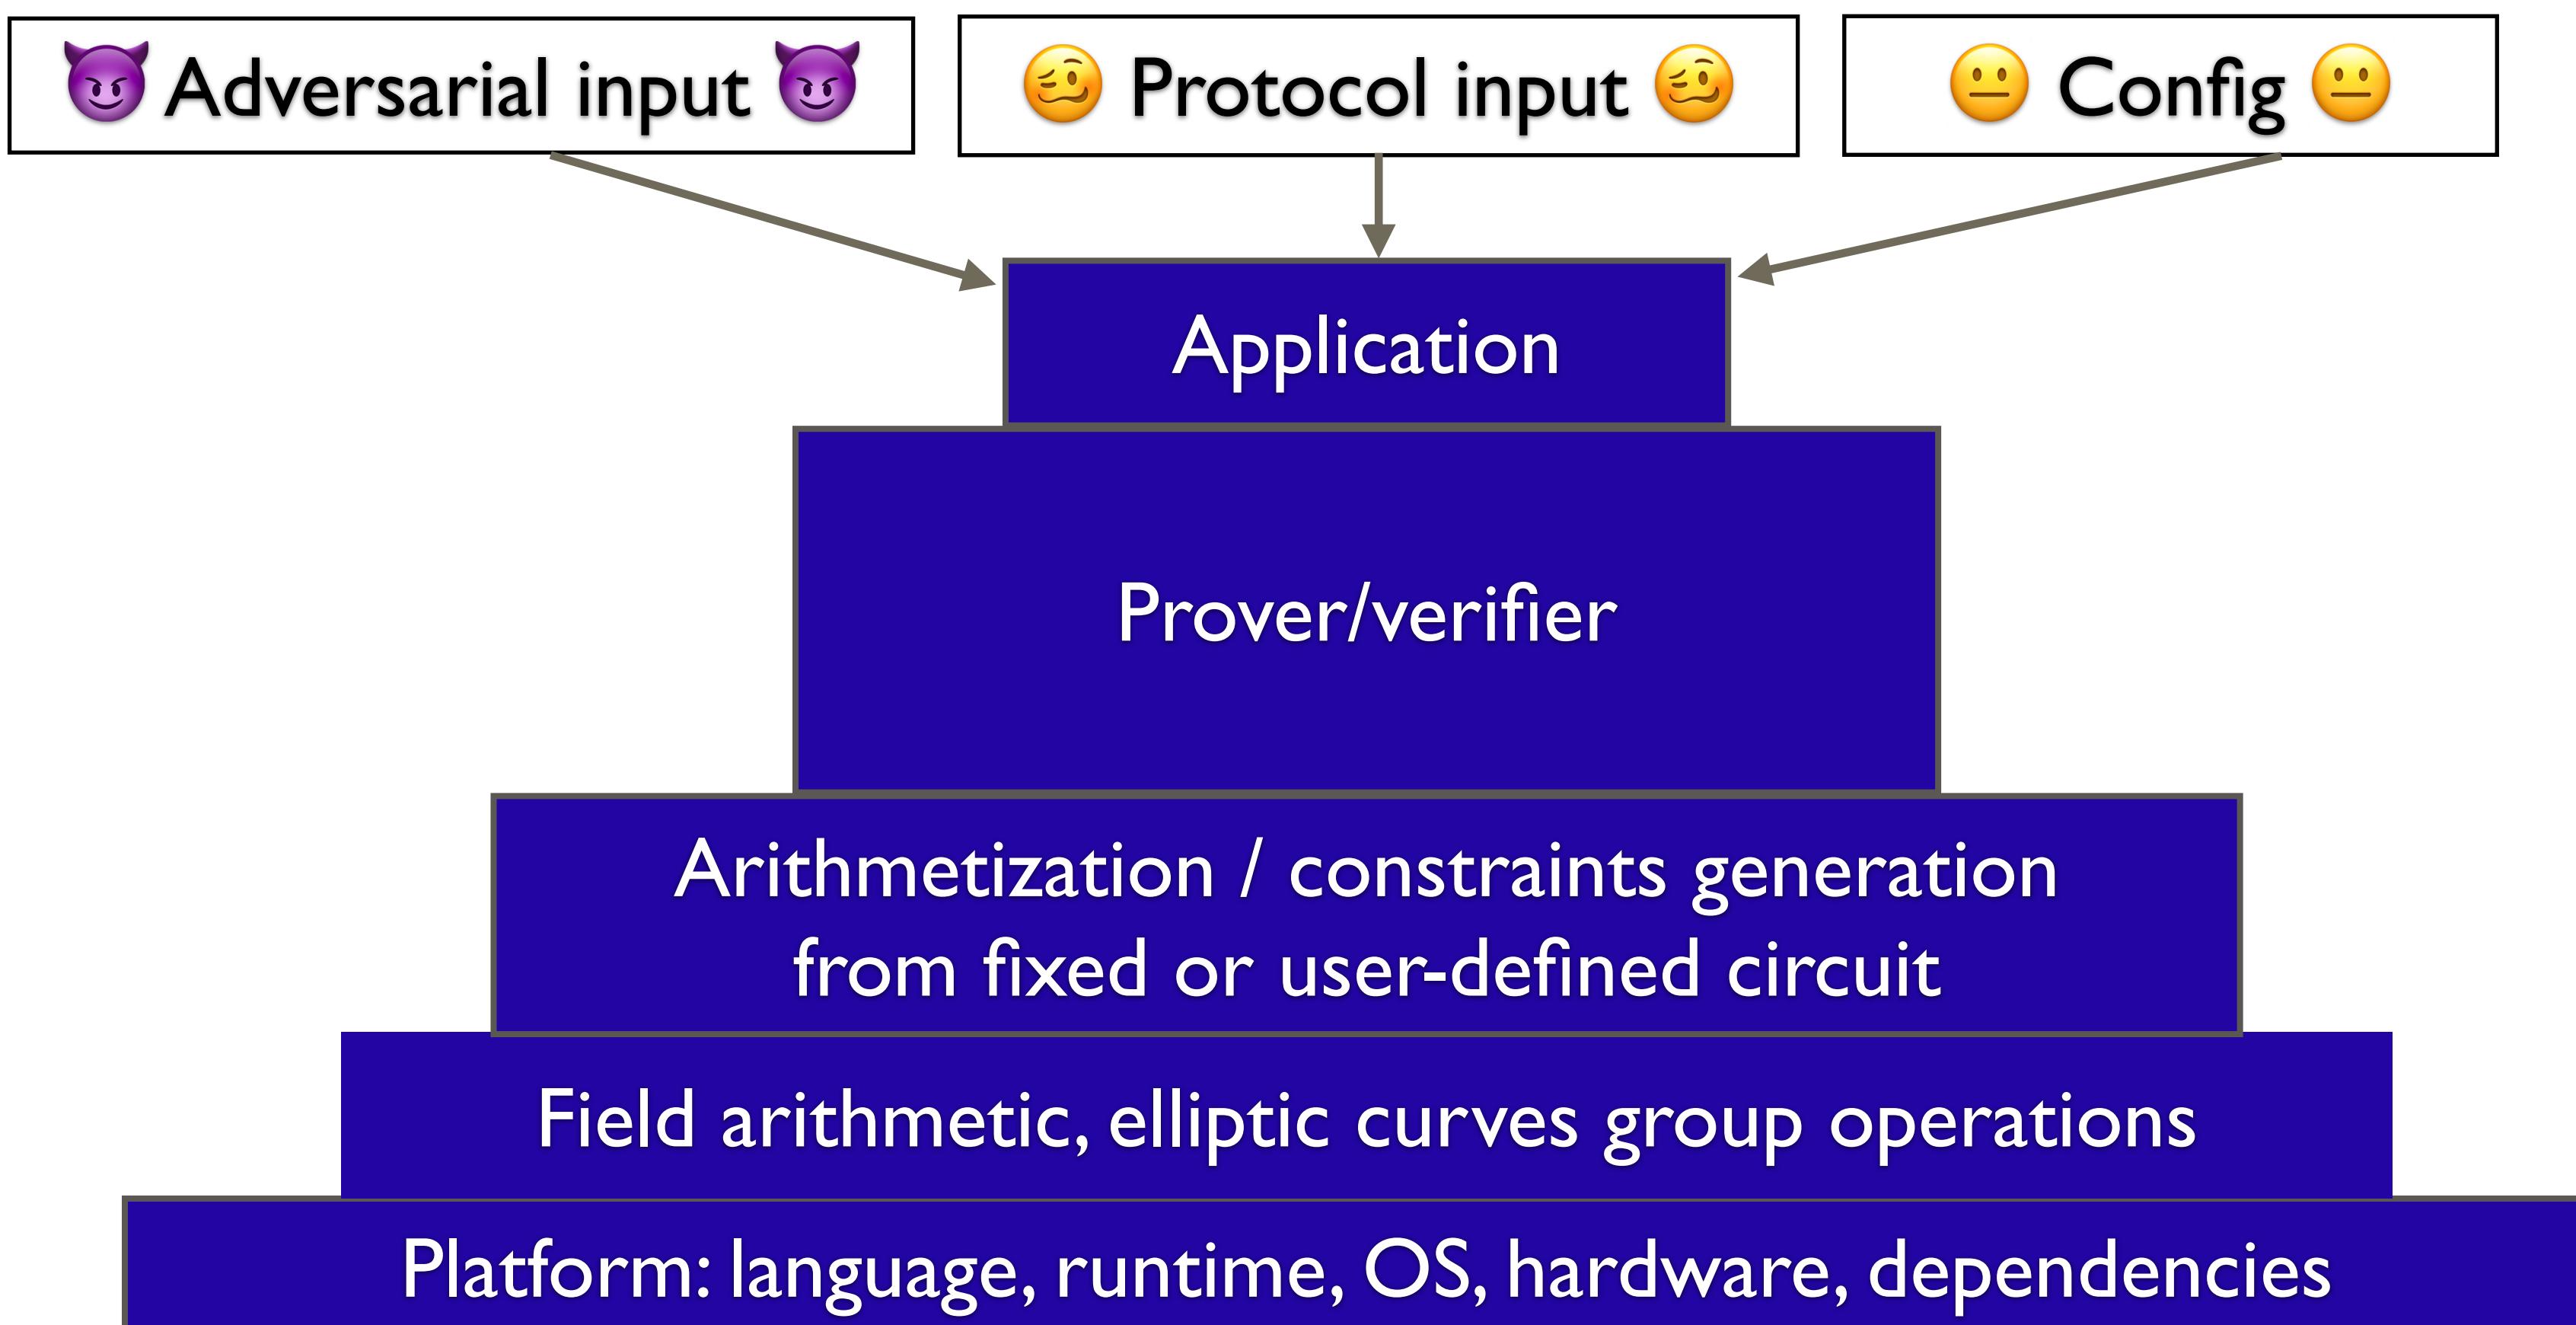
- Leakage of data, including via side channels, encodings (“ZK execution”)
- Any form in insecure state (code execution, DoS)

Compromise the supply-chain, via

- Trusted setup's code and execution
- Build and release process integrity
- Software dependencies
- **Break (on-chain) software** (incl. verifier) via smart contract bugs, logic flaws, etc.

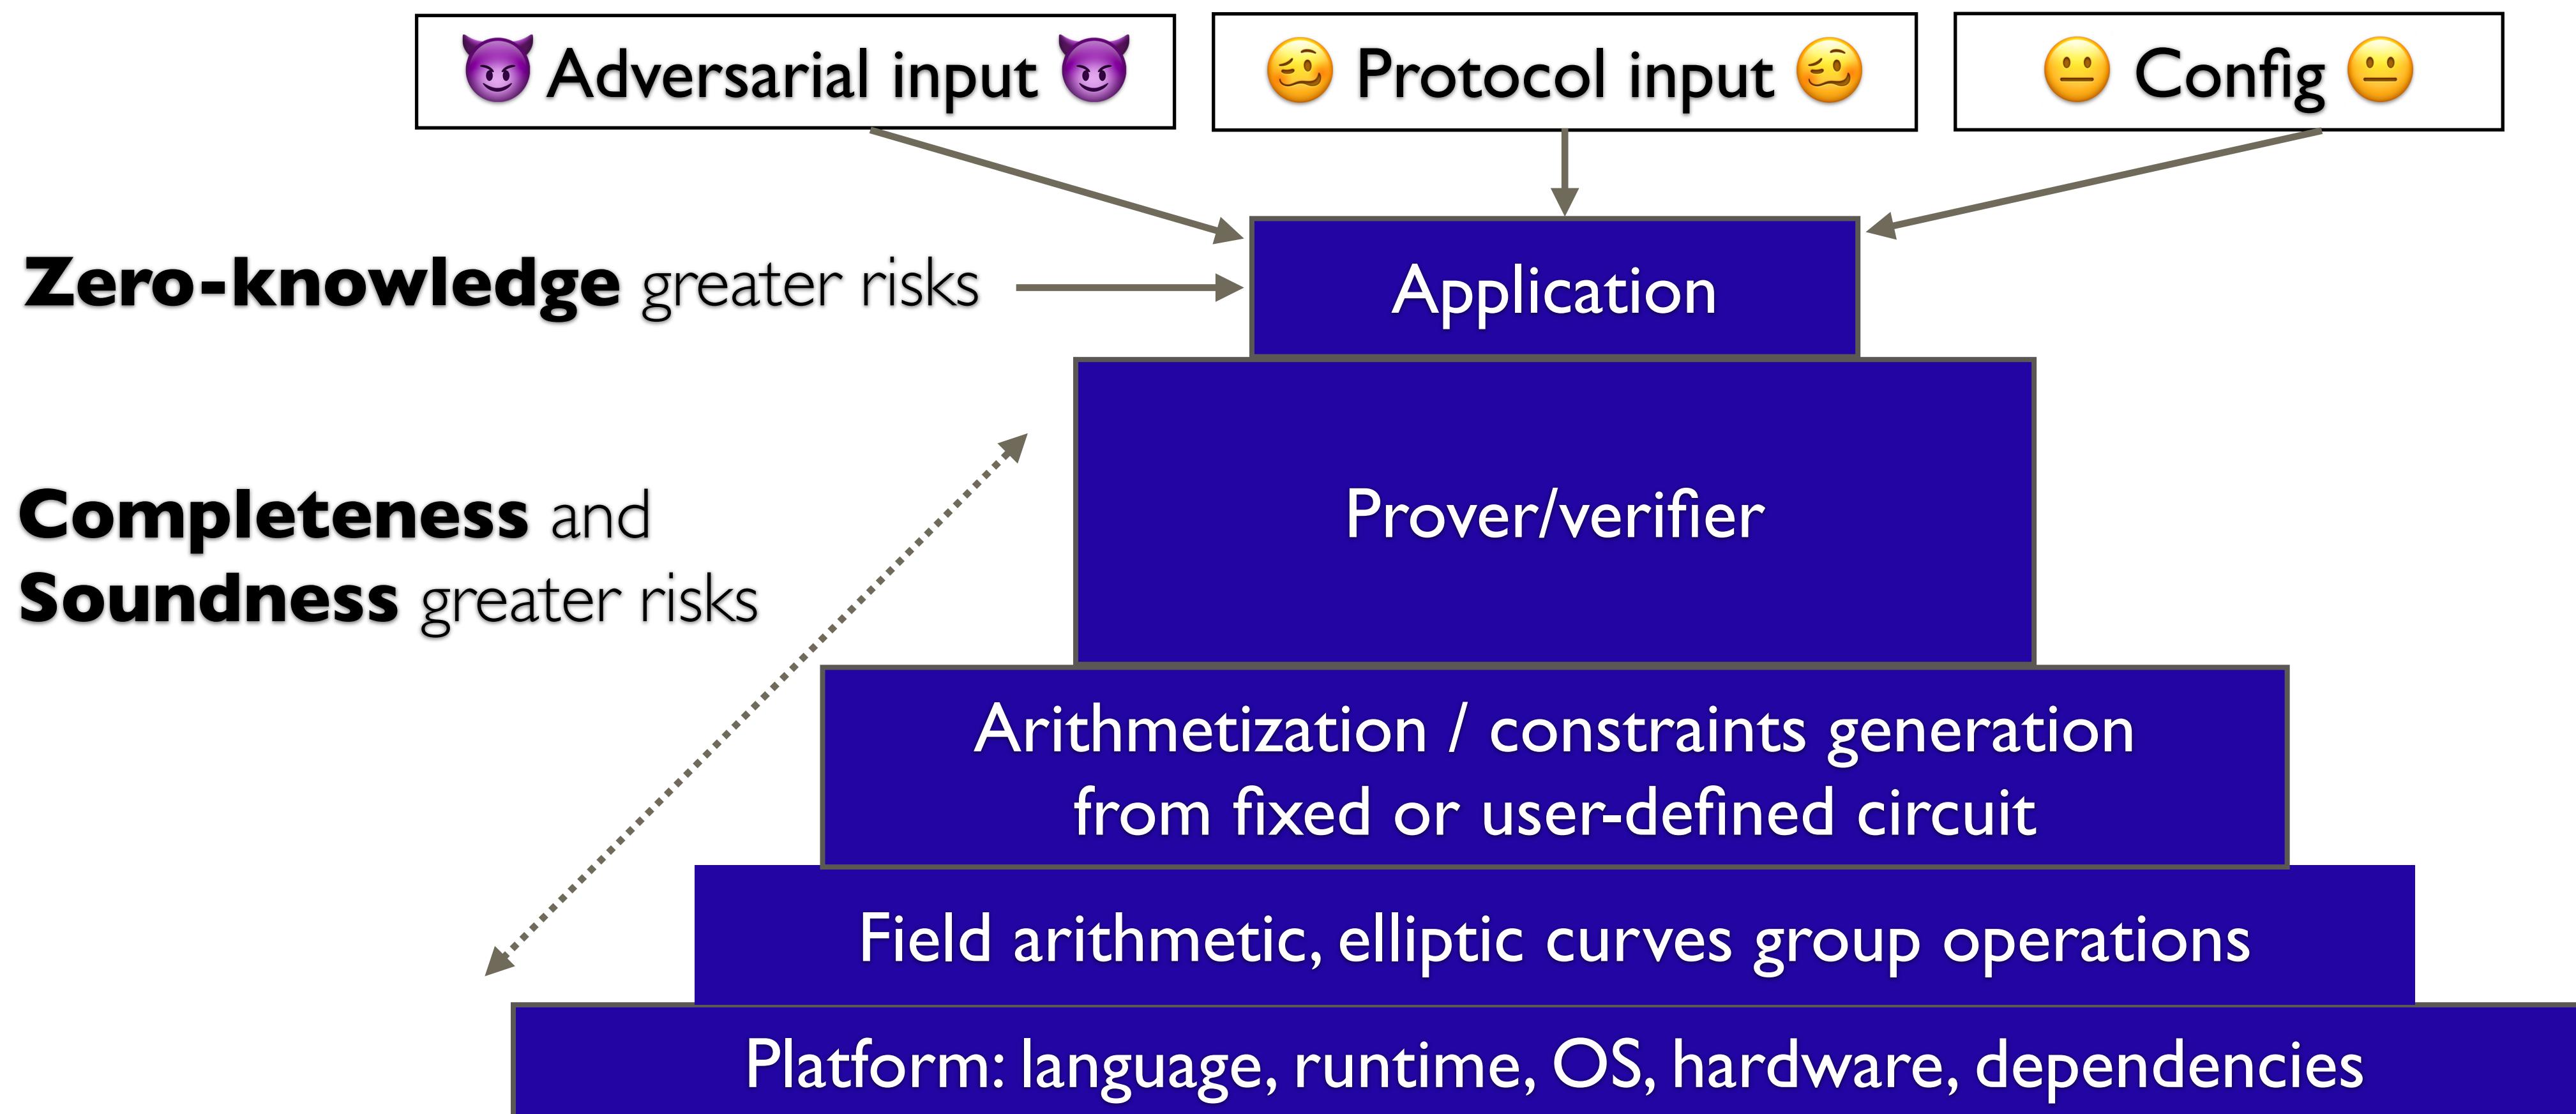
Multiple layers

- A failure in a **lower layer** can jeopardise the security of all upper layers



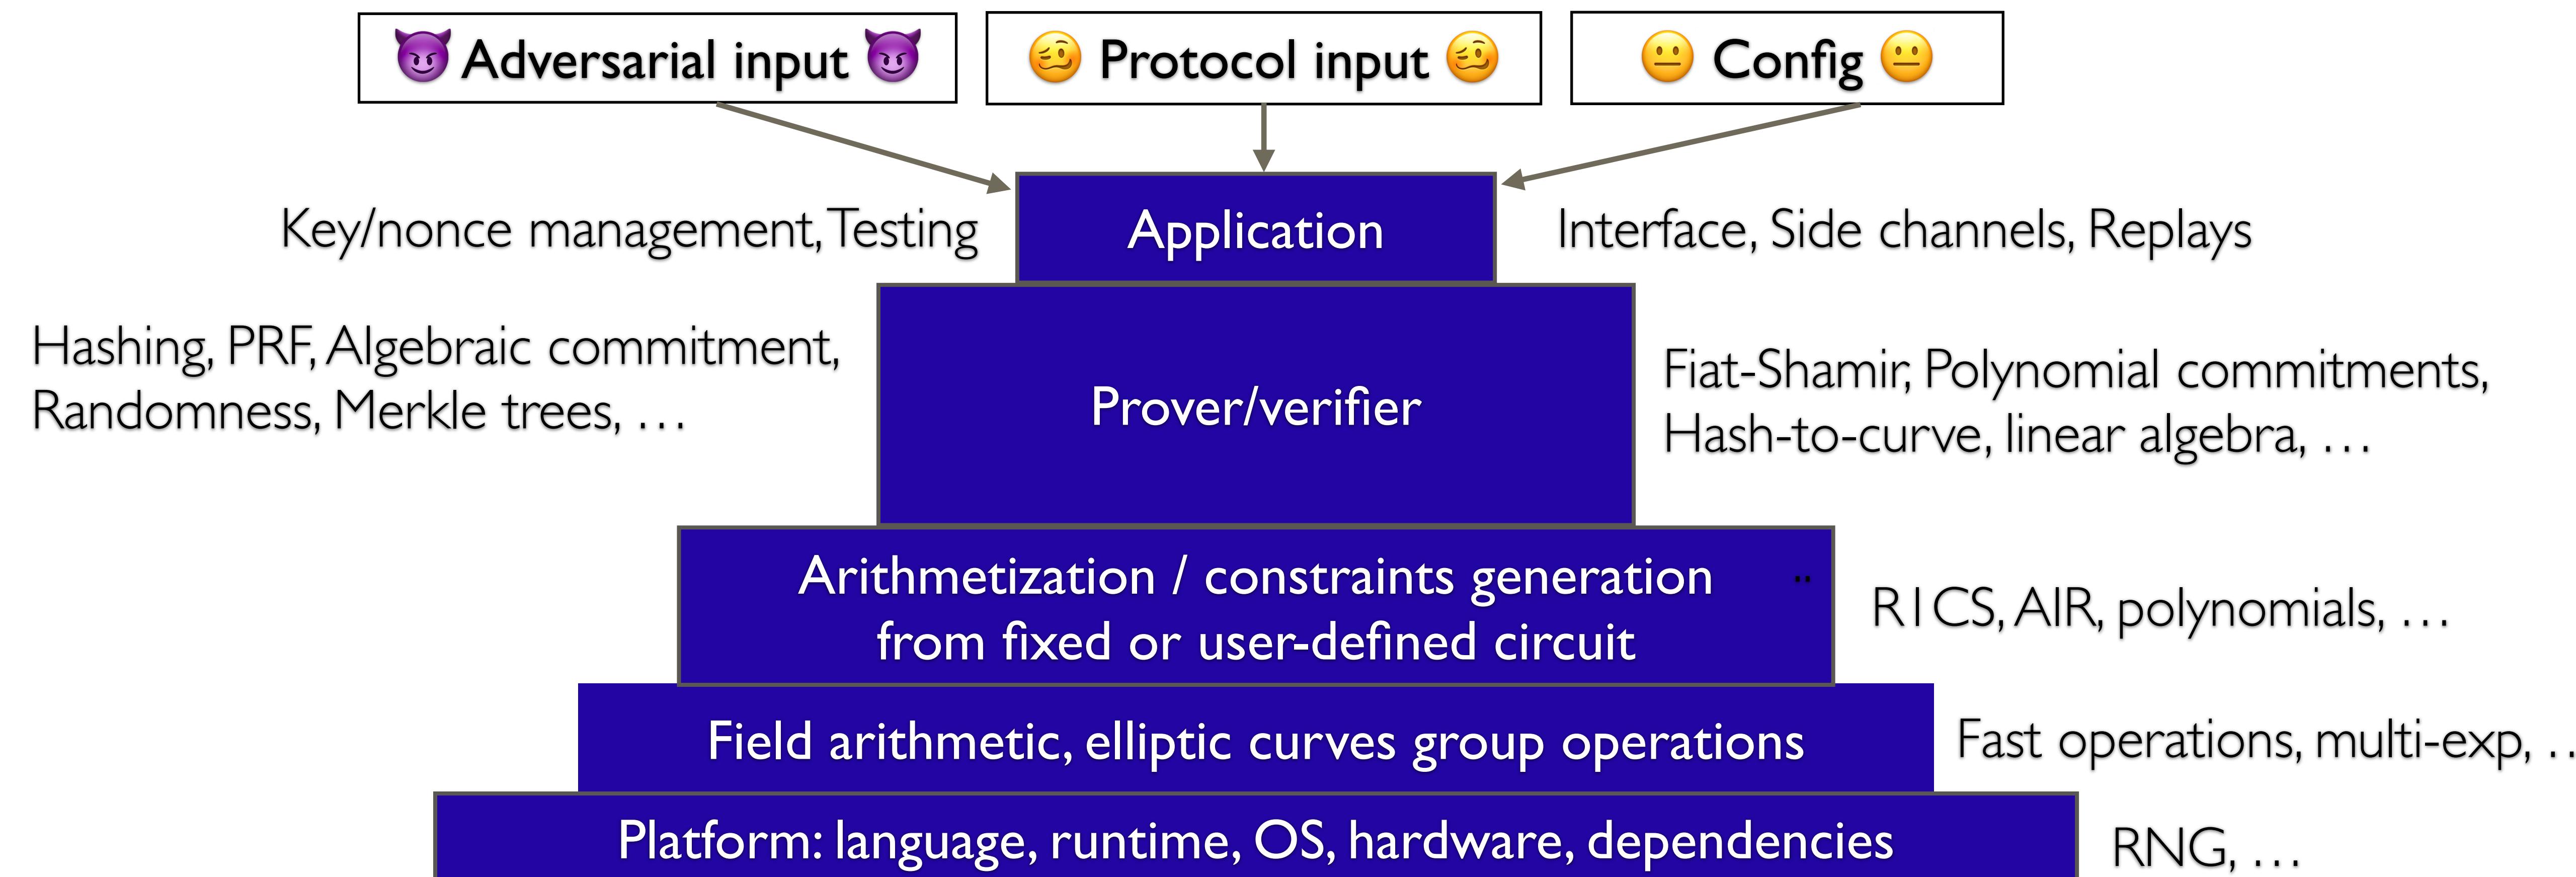
Multiple layers

- A failure in a **lower layer** can jeopardise the security of all upper layers



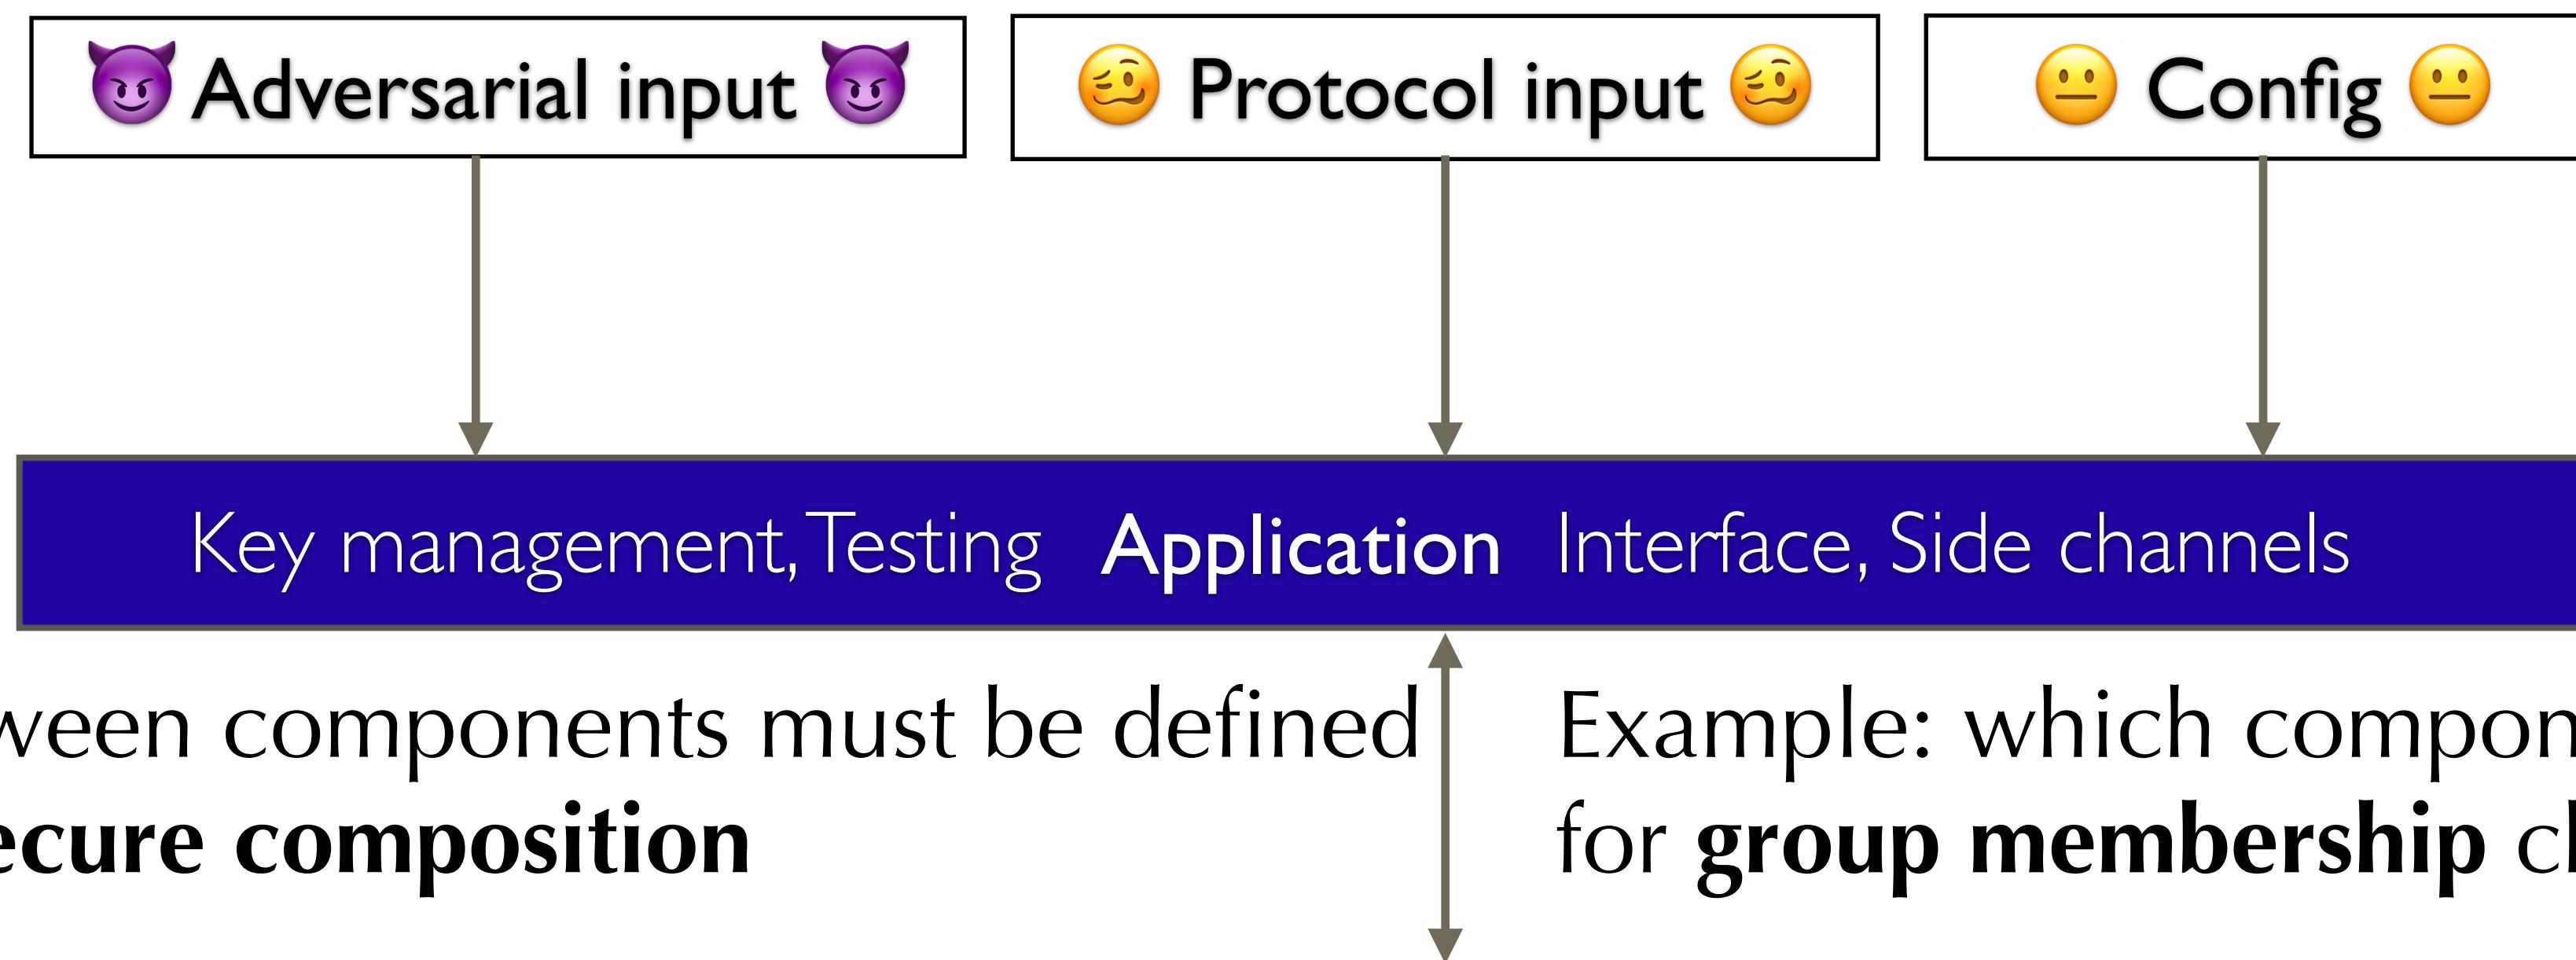
Multiple layers

- A failure in a **subcomponent** can jeopardise the security of all upper layers



Multiple layers

- Security 101: **Input validation** must be defined, implemented, and tested



Elliptic curves, Pairings, Hash functions, PRF, Algebraic commitment
Randomness, Merkle trees **Prover/verifier** Linear algebra, Multi-exp.
Polynomial commitments, Fiat-Shamir transforms, etc. etc.



Soundness – Field arithmetic

Vulnerability allowing double spend #16

 Closed poma opened this issue on 26 Jul 2019 · 2 comments



poma commented on 26 Jul 2019 · edited

...

Looks like in [Semaphore.sol#L83](#) we don't check that nullifier length is less than field modulus. So `nullifier_hash + 21888242871839275222246405745257275088548364400416034343698204186575808495617` will also pass snark proof verification if it fits into uint256, allowing double spend.

Root cause: Missing overflow check (of a nullifier ~ unique ID of a shielded payment)

<https://github.com/appliedzkp/semafore/issues/16>

Soundness – Field arithmetic

Potential security bug with the zk-SNARK verifier

Closed

weijiekoh opened this issue on 21 Mar 2020 · 2 comments · Fixed by #43



weijiekoh commented on 21 Mar 2020

...

Expected Behavior

The `Verifier.verify()` function, not the function that calls it (i.e.

`Shield.createMSA()` and `Shield.createP0()`, should require that each public input to the snark is less than the scalar field:

Missing overflow check (of a public input)

<https://github.com/eea-oasis/baseline/issues/34>

Soundness – Field arithmetic

```
diff --git a/Verifier.sol b/Verifier.sol
index 1a23f3d..a2a2a2c 100644
--- a/Verifier.sol
+++ b/Verifier.sol
@@ -219,7 +221,7 @@ contract Verifier {
    if (input.length + 1 != vk.IC.length) revert Pairing.InvalidProof();
    Pairing.G1Point memory vk_x = vk.IC[0];
    for (uint256 i = 0; i < input.length; i++) {
-        if (input[i] >= Pairing.SCALAR_MODULUS) revert Pairing.InvalidProof();
+        // By EIP196 the scalar_mul verifies it's input is in the correct range.
        vk_x = Pairing.addition(vk_x, Pairing.scalar_mul(vk.IC[i + 1], input[i]));
    }
}
```

Missing overflow check (of a public input)

<https://github.com/appliedzkp/seaphore/pull/96/>

Soundness – R1CS

Discuss: enforce `mul_by_inverse` #70

Merged weikengchen merged 7 commits into `master` from `fix-mul-by-inverse` on 6 Jul

Conversation 12 Commits 7 Checks 5 Files changed 3

weikengchen commented on 4 Jul 2021 · edited

Description

It seems that the `mul_by_inverse` implementation has a soundness issue that the newly allocated `d_inv` does not need to be the inverse of `d` but could be any value. This can be a soundness issue as the `poly` gadgets have used this API.

```
fn mul_by_inverse(&self, d: &Self) -> Result<Self, SynthesisError> {
    let d_inv = if self.is_constant() || d.is_constant() {
        d.inverse()?;
    } else {
        if self.is_constant() || d.is_constant() {
            let d_inv = d.inverse()?;
            Ok(d_inv * self)
        } else {
    }
```

RUSTSEC-2021-0075 History

Flaw in `FieldVar::mul_by_inverse` allows unsound R1CS constraint systems

Field element inverse property not enforced by the constraint system

<https://github.com/arkworks-rs/r1cs-std/pull/70>

Soundness – Hash validation

Technical Details

The bug was found by [Kobi Gurkan](#) in the zk-SNARK implementation of the [MIMC hash function in circomlib](#), that is used in Tornado for building the merkle tree of deposits. If everything works as expected, users prove that they have committed a leaf to that tree during deposit without revealing the commitment itself. The buggy version did not check that resulting MIMC hash is correct. The fix is simple: instead of using the `=` operator the `<==` operator should be used.

Coding error, allowing to fake the witness' Merkle root and forge proofs

<https://tornado-cash.medium.com/tornado-cash-got-hacked-by-us-b1e012a3c9a8>

Soundness – Paper / Setup

Background

On March 1, 2018, Ariel Gabizon, a cryptographer employed by the Zcash Company at the time, discovered a subtle cryptographic flaw in the [BCTV14] paper that describes the zk-SNARK construction used in the original launch of Zcash. The flaw allows an attacker to create counterfeit shielded value in any system that depends on parameters which are generated as described by the paper.

This vulnerability is so subtle that it evaded years of analysis by expert cryptographers focused on zero-knowledge proving systems and zk-SNARKs. In an analysis [Parno15] in 2015, Bryan Parno from Microsoft Research discovered a different mistake in the paper. However, the vulnerability we discovered appears to have evaded his analysis. The vulnerability also appears in the subversion zero-knowledge SNARK scheme of [Fuchsbauer17], where an adaptation of [BCTV14] inherits the flaw. The vulnerability also appears in the ADSNARK construction described in [BBFR14]. Finally, the vulnerability evaded the Zcash Company's own cryptography team, which includes experts in the field that had identified several flaws in other parts of the system.

Theoretical flaw in the paper's setup description (sensitive values not cleared)

<https://electriccoin.co/blog/zcash-counterfeiting-vulnerability-successfully-remediated/>

Zero-knowledge – Application (Aztec)

Issue #2

We discovered our method for removing spending keys involving an account nullifier, broke the sender privacy of transactions from the account; and consequently changed our key removal procedure to use an “Account Nonce” instead of the nullifier.

Issue #3

Our privacy circuit was not correctly including the user account nonce within the encrypted note cipher-text. This would have mean that a deprecated account, with an old nonce, would be able to spend any note owned by the account. We modified our circuit to include the account nonce when encrypting notes.

Nonces not correctly set by the application, breaking privacy

<https://medium.com/@jaosef/54dff729a24f> (Aztec 2.0 Pre-Launch Notes)

Zero-knowledge – Application (Zcash)

4.2 ITM Attack: Defeating *zk-SNARKs*

We can think of this attack as a "defeat" of zero-knowledge mathematics only in practice, not in theory. Many qualifications are needed. We in no way "broke" the mathematics of *zk-SNARKs*, we are taking advantage of how *zk-SNARKs* are being used in higher level protocols, i.e. the Zcash Transaction Format Protocol and its associated consensus rules.

So *zk-SNARKs* are sound and we have not actually leaked knowledge directly from a zero-knowledge proof, that is mathematically impossible. We have leaked knowledge from how these proofs are used in the larger system called Zcash Protocol, itself an extension of Bitcoin Protocol which notoriously leaks metadata.

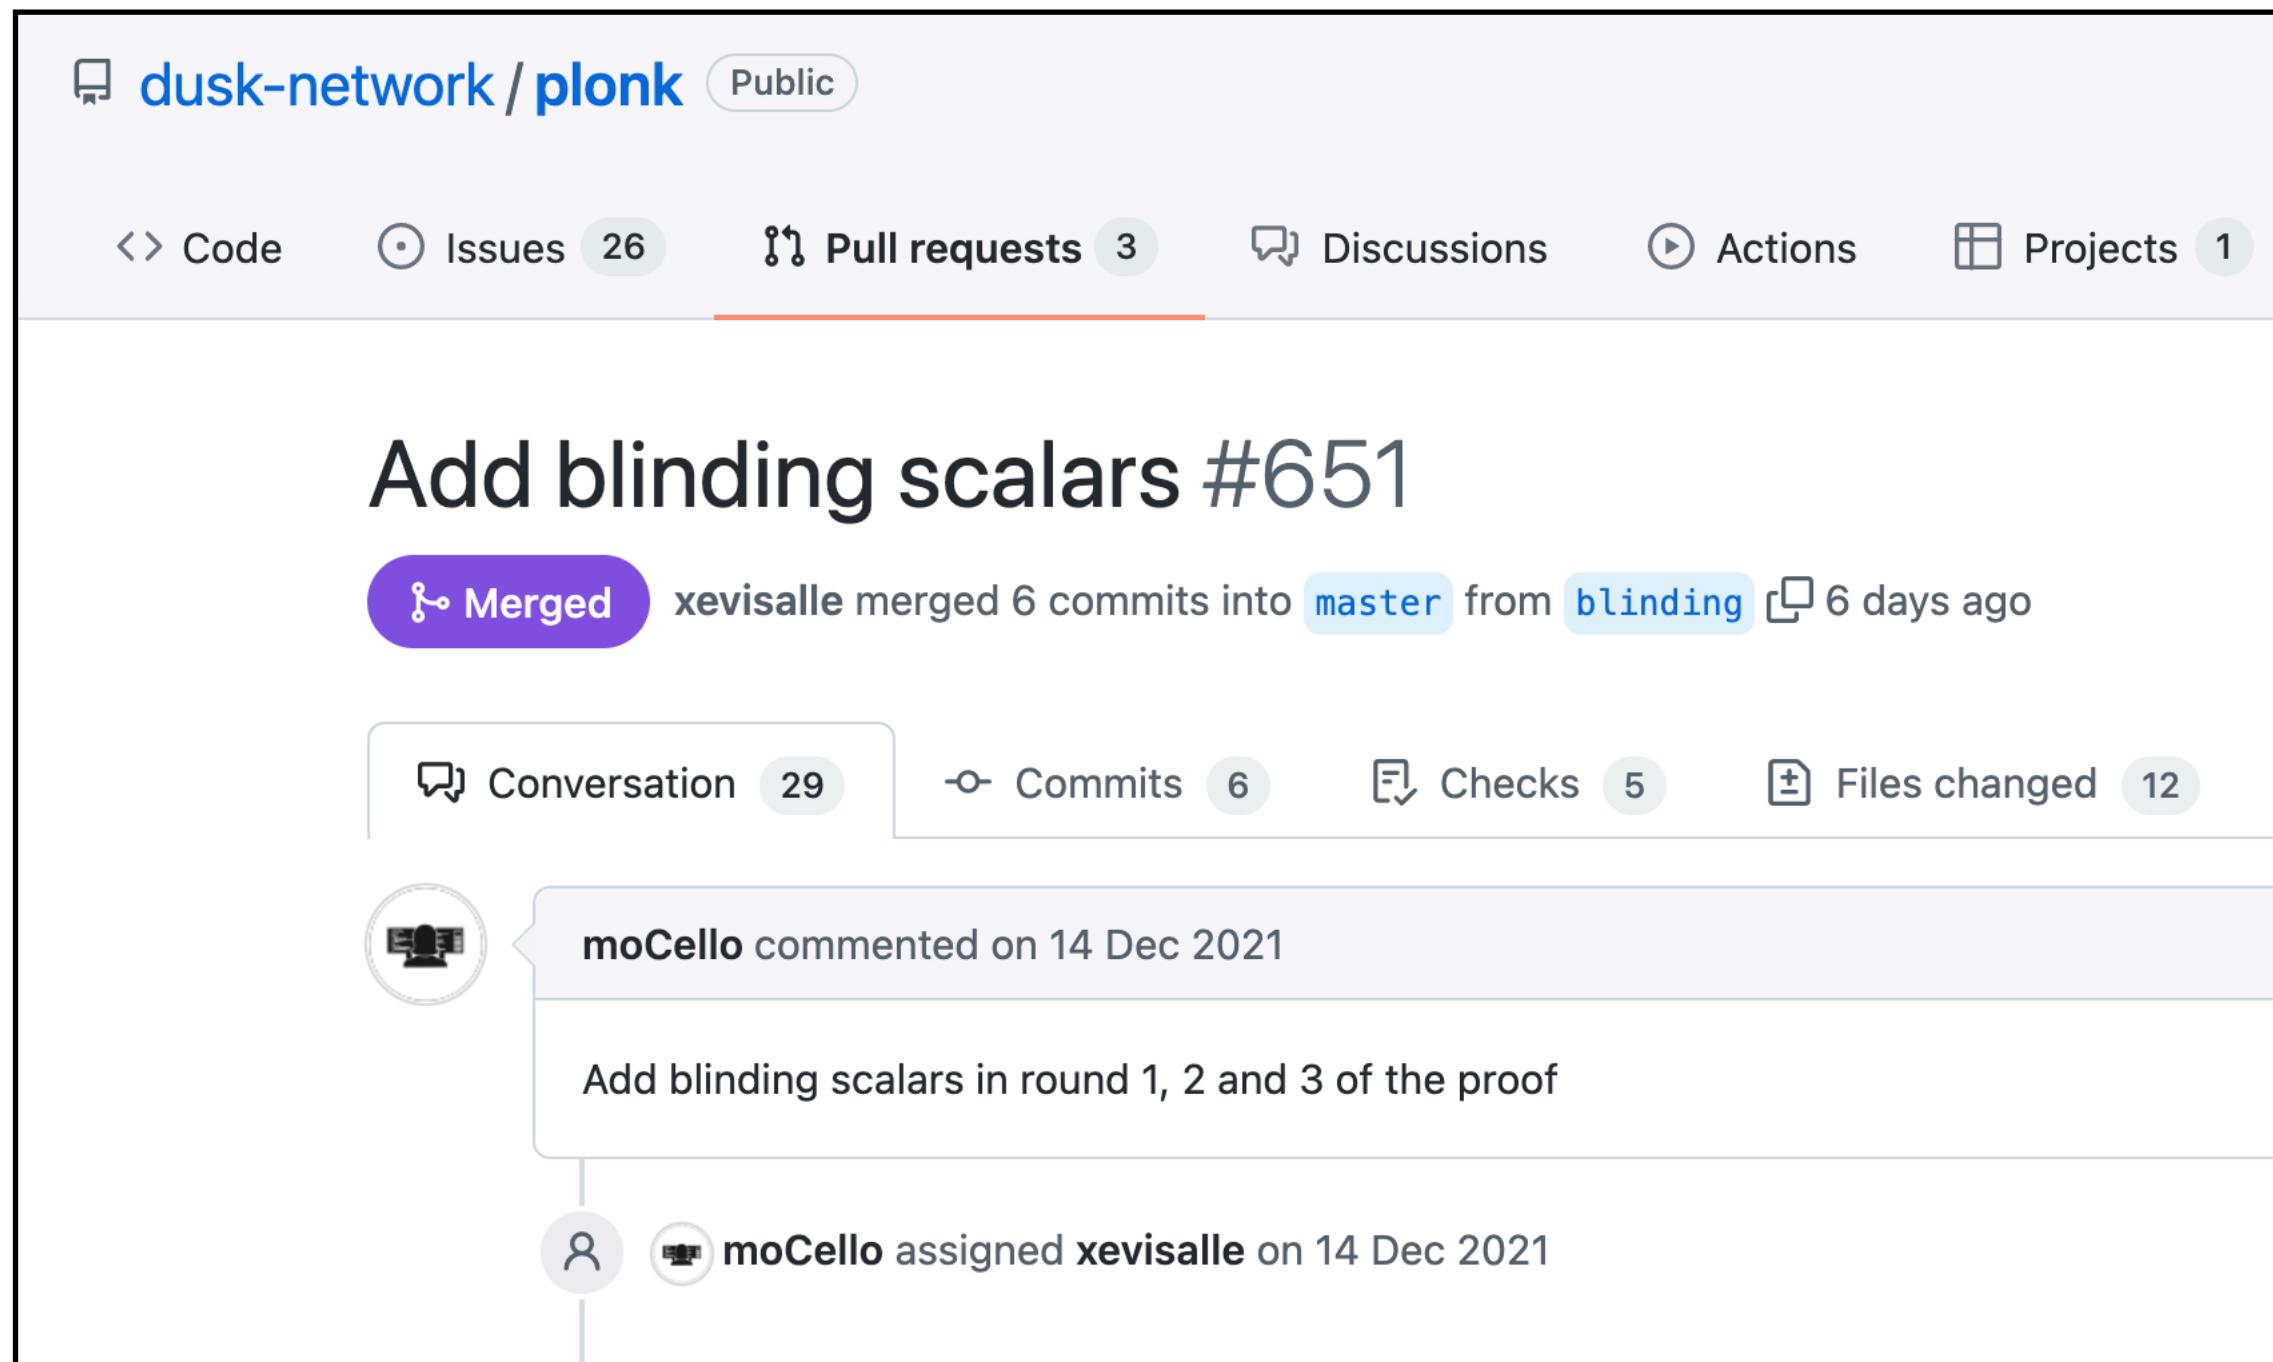
Correlations between (shielded) transactions leaking exploitable information

<https://eprint.iacr.org/2020/627.pdf>

Attacking Zcash Protocol For Fun And Profit
Whitepaper Version 0.5

Duke Leto + The Hush Developers[†]

Zero-knowledge – Prover (Plonkup)



Missing (randomized) blinding to hide private inputs – *potential ZK loss*

<https://github.com/dusk-network/plonk/pull/651>

DoS – Merkle tree

Bug: Merkle root position check

The rollup contains a “root” tree; a Merkle tree containing the past Merkle roots of the note tree (which contains all join-split “value” notes and user “account” notes).

As part of the root rollup circuit, the rollup provider must compute the new root of the note tree and insert it into the root tree. The *intended* position of the new leaf in the tree is directly adjacent to the rightmost non-zero leaf; i.e. the tree is initialized to all zero leaves, and then updated from left-to-right.

The bug was that our circuit did not actually constrain the position of the new leaf. In reality, the rollup provider could have inserted the new leaf at **any** position in the root tree. An adversary would have been able to insert a leaf at an arbitrary location in the root tree *and not reveal the location (this location is not a public input)*.

If after such an insertion the adversary doesn’t participate in future rollup creation, *from that point on *nobody else** can create a valid rollup and the system is frozen and unable to process for any future transactions.

Incomplete tree constraints, leading to a freeze of the rollup validation

<https://medium.com/aztec-protocol/vulnerabilities-found-in-aztec-2-0-9b80c8bf416c>

DoS / Completeness? – DSL / Signatures

veorq commented yesterday

`crypto.signature.signature.verify()` rejects signatures with an `r`, inverse `s`, or message (hash) greater than `2**251 < EC_ORDER`:

[cairo-lang/src/starkware/crypto/starkware/crypto/signature/signature.py](#)
Lines 199 to 201 in 4e23351

```
199     assert 1 <= r < 2 ** N_ELEMENT_BITS_ECDSA, "r = %s" % r
200     assert 1 <= w < 2 ** N_ELEMENT_BITS_ECDSA, "w = %s" % w
201     assert 0 <= msg_hash < 2 ** N_ELEMENT_BITS_ECDSA, "msg_hash = %s" % msg_hash
```

There's a gap of $\sim 2^{196}$ values, thus a probability to hit an invalid `r` or `s` that is of the order of $2^{(196-251)/2} = 2^{54}$, when generating an ECDSA sig for some fixed message using a standard algorithm (rather than Cairo's `sign()`, which enforces these constraints).

I can't think of a specific attack scenario at the moment, but I would expect to find applications where either

1. that accidental failure rate would be unacceptably high, or
2. adversaries could bruteforce invalid sigs to do some kind of DoS, or worse (with plausible deniability)

I probably miss some of the context, and you may have a good reason to verify sigs that way.

Valid signatures rejected, risk initially deemed negligible

<https://github.com/starkware-labs/cairo-lang/issues/39>

Other types of bugs

- **Crypto** issues such as:
 - Pedersen bases generation/uniqueness
 - Padding scheme in algebraic hashes and commitments
 - Non-conform implementations of crypto schemes (e.g. Poseidon algebra bugs)
 - Insufficient data being “Fiat-Shamir'd” from the transcript
- **Composability**: unsafe interactions between nested proof systems
- **Side channels**: Non-ct code, RAM leakage, speculative execution leaks

Conclusions

😊 Why not be too scared?

- Robust code and frameworks (e.g. Rust projects such as arkworks and zkcrypto)
- DSLs (Cairo, Leo, etc.) make it easier to write safe code
- Relatively narrow attack surface in practice

Conclusions

😊 Why not be too scared?

- Robust code and frameworks (e.g. Rust projects such as arkworks and zkcrypto)
- DSLs (Cairo, Leo, etc.) make it easier to write safe code
- Relatively narrow attack surface in practice

😱 Why be scared?

- Few people understand zkSNARKs, even fewer can find bugs
- Lack of tooling (wrt testing, fuzzing, verification)
- More ZKPs used => more \$\$\$ at stake => greater ROI for vuln researchers

Thank you!



JP Aumasson
@veorq

CSO @ taurushq.com

Thanks to Aleo, Protocol Labs, Kobi Gurkan, Adrian Hamelink, Lúcás Meier, Mathilde Raynal